

HPCL Ass.4 :

BE_34_Samruddhi Khairnar

15/04

Write a CUDA Program for :

1. Addition of two large vectors
2. Matrix Multiplication using CUDA C.

Code :

```
// HPCL_4_BE_34 - Vector addition, Matrix Multiplication - CUDA C
```

```
# include <stdio.h>
```

```
__global__ void add_vectors(int *g_a, int *g_b, int *g_c)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    g_c[i] = g_a[i] + g_b[i];
}
```

```
void run_vector_addition()
{
    printf("Enter the length of both the vectors : ");
    int n;
    scanf("%d", &n);
    int c_a[n], c_b[n], c_c[n];
    for (int i=0;i<n;i++)
    {
        printf("Enter element %d of Vector 1 : ", i);
        scanf("%d", &c_a[i]);
        printf("Enter element %d of Vector 2 : ", i);
        scanf("%d", &c_b[i]);
    }
}
```

```
int *g_a, *g_b, *g_c;
```

```
int size = n*sizeof(int);
cudaMalloc(&g_a, size);
cudaMalloc(&g_b, size);
cudaMalloc(&g_c, size);
cudaMemcpy(g_a, c_a, size, cudaMemcpyHostToDevice);
cudaMemcpy(g_b, c_b, size, cudaMemcpyHostToDevice);
cudaMemcpy(g_c, c_c, size, cudaMemcpyHostToDevice);
```

```
add_vectors<<<5,5>>>(g_a, g_b, g_c); // 5 Blocks * 5 Threads = 25
Threads in total.
```

```
cudaDeviceSynchronize();
cudaMemcpy(c_c, g_c, size, cudaMemcpyDeviceToHost);

cudaFree(g_a);
cudaFree(g_b);
cudaFree(g_c);

printf(" ");
for (int i=0;i<n;i++)
    printf("%d ", c_a[i]);
printf("\n+");
for (int i=0;i<n;i++)
    printf("%d ", c_b[i]);
printf("\n=");
for (int i=0;i<n;i++)
    printf("%d ", c_c[i]);
printf("\n");
}

__global__ void matrix_multiply(int *g_a, int *g_b, int *g_c, int size)
{

/*
```

*I have converted the 2D array of 3*3 elements into a 1D array of 9 elements hence multiplying the block dimension to the row index of each of the matrices - A, B, C. Shinde ma'am has used 2D array hence x,y dimensions.*

	0	1	2	
0	{A00,A01,A02}	{B00,B01,B02}	{C00,C01,C02}	
1	{A10,A11,A12}	X {B10,B11,B12}	= {C10,C11,C12}	
2	{A20,A21,A22}	{B20,B21,B22}	{C20,C21,C22}	

$C00 = A00*B00 + A01*B10 + A02*B20$
 $C01 = A00*B01 + A01*B11 + A02*B21$

Hence, we must have a loop - from 0 to 2:
 $C00, C01 = 0, 0$

```
for i in range(0,3):
    C00 += A0i * Bi0
    C01 += A0i * Bi1
```

Hence, in a loop, $C[Block_id][thread_id] += A[Block_id][i] *$

```
B[i][thread_id]
```

And as I have converted all of them into 1D arrays, The left index is multiplied by Block_dim (no.of threads in a block)

```
for i in range(0,3):
```

```
    C[Block_id*Block_dim + thread_id] += A[Block_id*Block_dim + i] *  
    B[i*Block_dim +thread_id]
```

```
*/
```

```
    int idx = blockDim.x * blockIdx.x + threadIdx.x;  
    for(int i=0; i<size; i++)  
        g_c[idx] += g_a[blockDim.x * blockIdx.x + i] * g_b[i * blockDim.x +  
threadIdx.x];  
}
```

```
void run_matrix_multiplication()
```

```
{  
    int n=3;  
    int c_c[3][3] = {{0,0,0},  
                     {0,0,0},  
                     {0,0,0}};  
  
    int c_a[n][n], c_b[n][n];  
    for (int i=0;i<n;i++)  
    {  
        for (int j=0;j<n;j++)  
        {  
            printf("Enter element %d,%d of Matrix A : ", i, j);  
            scanf("%d", &c_a[i][j]);  
            printf("Enter element %d,%d of Matrix B : ", i, j);  
            scanf("%d", &c_b[i][j]);  
        }  
    }  
}
```

```
int *g_a, *g_b, *g_c;
```

```
int size = n*n*sizeof(int);  
cudaMalloc(&g_a, size);  
cudaMalloc(&g_b, size);  
cudaMalloc(&g_c, size);  
cudaMemcpy(g_a, c_a, size, cudaMemcpyHostToDevice);  
cudaMemcpy(g_b, c_b, size, cudaMemcpyHostToDevice);  
cudaMemcpy(g_c, c_c, size, cudaMemcpyHostToDevice);
```

```
matrix_multiply<<<3,3>>>(g_a, g_b, g_c, size); // 3 Blocks, each having
```

3 threads

```
cudaDeviceSynchronize();
cudaMemcpy(c_c, g_c, size, cudaMemcpyDeviceToHost);

cudaFree(g_a);
cudaFree(g_b);
cudaFree(g_c);

printf("A = \n");
for (int i=0;i<n;i++)
{
    for (int j=0;j<n;j++)
        printf("%d,", c_a[i][j]);
    printf("\n\n");
}
printf("B = \n");
for (int i=0;i<n;i++)
{
    for (int j=0;j<n;j++)
        printf("%d,", c_b[i][j]);
    printf("\n\n");
}
printf("Multiplication = \n");
for (int i=0;i<n;i++)
{
    for (int j=0;j<n;j++)
        printf("%d,", c_c[i][j]);
    printf("\n\n");
}
}

int main()
{
    int ch;
    while(true)
    {
        printf("Enter 1-Vector Addition | 2-Matrix Multiplication | 0-Exit\n");
        scanf("%d", &ch);
        if(ch==1)
            run_vector_addition();
        else if(ch==2)
            run_matrix_multiplication();
        else if(ch==0)
        {
            printf("Exited Successfully.\n");
        }
    }
}
```

```

        break;
    }
    else
        printf("Invalid input.\n");
}

return 0;
}

```

Output :

```

comp-proj-sys05@compprojsys05-OptiPlex-3010:~/Downloads$ nvcc HPCL_4_BE_34.cu
comp-proj-sys05@compprojsys05-OptiPlex-3010:~/Downloads$ ./a.out
Enter 1-Vector Addition | 2-Matrix Multiplication | 0-Exit : 1
Enter the length of both the vectors : 5
Enter element 0 of Vector 1 : 1
Enter element 0 of Vector 2 : 1
Enter element 1 of Vector 1 : 2
Enter element 1 of Vector 2 : 2
Enter element 2 of Vector 1 : 3
Enter element 2 of Vector 2 : 3
Enter element 3 of Vector 1 : 4
Enter element 3 of Vector 2 : 4
Enter element 4 of Vector 1 : 2
Enter element 4 of Vector 2 : 2
  1 2 3 4 2
+1 2 3 4 2
=2 4 6 8 4
Enter 1-Vector Addition | 2-Matrix Multiplication | 0-Exit : 2
Enter element 0,0 of Matrix A : 1
Enter element 0,0 of Matrix B : 1
Enter element 0,1 of Matrix A : 2
Enter element 0,1 of Matrix B : 0
Enter element 0,2 of Matrix A : 3
Enter element 0,2 of Matrix B : 0
Enter element 1,0 of Matrix A : 4
Enter element 1,0 of Matrix B : 0
Enter element 1,1 of Matrix A : 5
Enter element 1,1 of Matrix B : 1
Enter element 1,2 of Matrix A : 6
Enter element 1,2 of Matrix B : 0
Enter element 2,0 of Matrix A : 7
Enter element 2,0 of Matrix B : 0
Enter element 2,1 of Matrix A : 8
Enter element 2,1 of Matrix B : 0

```

Enter element 2,2 of Matrix A : 9

Enter element 2,2 of Matrix B : 1

A =

1,2,3,

4,5,6,

7,8,9,

B =

1,0,0,

0,1,0,

0,0,1,

Multiplication =

1,2,3,

4,5,6,

7,8,9,

Enter 1-Vector Addition | 2-Matrix Multiplication | 0-Exit : 0

Exited Successfully.